

The Use of Quantum Support Vector Machines for Stock Price Prediction

Ben Muir and Hasan Fakhri

April 8th 2025

1 Introduction

1.1 Context

Quantum computing has emerged as a promising tool for solving computationally intensive problems, including those found in financial modeling. In recent years, financial institutions have increasingly turned to Artificial Intelligence (AI) and Machine Learning (ML) techniques to perform predictive analytics on stock market data [1]. Among classical ML models, Support Vector Machine (SVM)s have demonstrated solid performance for classification tasks such as stock movement direction [2]. However, SVMs encounter limitations when dealing with high-dimensional or non-linearly separable data.

Quantum Support Vector Machine (QSVM)s, a quantum version of classical SVMs, exploit the properties of quantum mechanics like superposition and entanglement to enhance classification capabilities and computational efficiency. These models leverage quantum feature spaces and quantum kernels to transform classical datasets into forms that can be more efficiently classified on quantum hardware or simulators [3]. Furthermore, research has shown that quantum-enhanced models, particularly QSVMs, have the potential to outperform classical models in specific financial contexts, particularly when paired with dimensionality reduction and robust feature engineering [6].

1.2 Problem Statement

Stock price prediction remains one of the most complex challenges in financial analytics due to the highly volatile nature of stock market data. Classical ML models struggle to maintain accuracy when scaling across large datasets with multiple correlated indicators. The primary objective of this project is to evaluate the effectiveness of QSVMs for binary classification of stock price movement, using real market data and derived indicators. A secondary objective is to test whether the quantum model can achieve competitive or superior performance compared to existing models.

1.3 Result

This project successfully implements a QSVM-based classification system using Qiskit’s quantum machine learning library. The system is tested across several feature combinations and evaluated on real historical stock data from companies like Apple, Microsoft, Visa, and Honeywell. The final model achieves an accuracy of 59.4% and an F1-score of 67.5% using a two-feature quantum kernel derived from technical indicators (Moving Average Convergence Divergence (MACD) and Exponential Moving Average (EMA) ratio), surpassing the benchmark results reported in prior work using classical or Principal Component Analysis (PCA)-enhanced quantum models [6].

1.4 Outline

The remainder of this report is structured as follows. Section 2 presents background information relevant to quantum computing and financial ML. Section 3 describes in detail the experimental methodology and implementation. The evaluation of the model is discussed in Section 4. We conclude with Section 5.

2 Background Information

Quantum Machine Learning (QML) is a growing field that combines the computational advantages of quantum algorithms with the learning capabilities of classical models. Among these, the QSVM model has shown particular promise in binary classification tasks, offering potential advantages over classical SVMs in speed and dimensionality handling [4].

Classical approaches to stock prediction typically involve techniques such as (Artificial Neural Network (ANN)), (Long Short-Term Memory (LSTM)) models, and SVMs [2]. These models rely heavily on carefully engineered features derived from historical stock data, such as momentum, trading volume, and trend-based indicators like the MACD or Relative Strength Index (RSI). However, they often struggle with generalization and suffer from overfitting, especially in high-dimensional feature spaces.

Quantum models, by contrast, can encode complex feature maps using quantum circuits like the ZZFeatureMap in Qiskit, allowing transformation of non-linearly separable data into higher-dimensional Hilbert spaces. This enables better class separation using quantum kernels such as the FidelityQuantumKernel, as used in this project [3].

Recent studies have investigated the application of QSVMs for financial tasks. In particular, Srivastava et al. (2023) applied Quantum Annealing and PCA for feature selection before using a QSVM to classify stock price movements. Their results suggest that careful dimensionality reduction can enhance model performance, but that a well-chosen set of stock indicators can often yield comparable or superior results without PCA [6]. Similarly, Bhasin et al. (2024) demonstrated improved portfolio management results using optimized QSVM implementations [4].

In the current project, we apply these principles to design a QSVM framework that can efficiently train on a subset of selected stock features and produce accurate directional predictions.

2.1 Data and Processing

2.1.1 Dataset and Pre-treating

The dataset we use in this project originated from Kaggle [5]. The dataset provided us with seven different metrics to base our features upon. These metrics included: opening price, closing price, high, low, adjusted-closing price, and volume. The dataset is quite large, as it includes all stock tickers currently trading on the NASDAQ, therefore we picked out some example stock tickers to use (Nvidia, Acilent, Microsoft) as well as the stock tickers used in the study we are comparing with (Honeywell, Visa, Apple, Johnson and Johnson). The mix of stock tickers allows for the QSVM to train on multiple different types of stocks. In order for the QSVM to properly learn from the data it must be pre-treated to ensure numerical stability and decrease the risk of noise. This project used the same method of Min-Max normalization with a scale from 0 to 1, as done in this study [4].

2.1.2 Stock Indicators

Certain stock indicators were replicated from the study to ensure we are consistent with our comparisons and to determine optimal predictions, these include the (MACD), (EMA-12 and 26), (Average True Range (ATR)), and Relative Strength Index (RSI). Two simplistic indicators were also calculated based on the technical indicators of the tickers, Price-Volume Ratio and Momentum. These indicators were calculated within the software by using builtin functions (such as for EMA), helper functions (such as for ATR and RSI), or by calculating the indicator using the formula and the values within the dataset (such as for Momentum or Previous Close).

2.2 Feature Selection

The paper that our project is based upon utilizes a Quantum Annealing algorithm to develop feature selection as well as dimensionality reduction prior to the training and predictions [6]. Initially we planned to have a Quantum Annealing algorithm using a Quadratic Unconstrained Binary Optimization (QUBO), this algorithm would be able to calculate the optimal features to be selected for the QSVM. However the algorithm's selection did not result in high enough metrics to be determined as useful. The decision was made to select the features manually by testing different combinations of both raw input data from the dataset and calculated stock indicators which are present in the study [6].

Prior to any implementation of the QSVM, common stock metrics for the dataset were created to better develop plans towards an optimal feature selection process. The volatility of a stock is an extremely important factor for determining the optimal feature to be selected. We can see in figure 1 that the stock involved in the tech industry (AAPL, NVDA, MSFT) have higher volatility numbers, which means these stocks could be predicted much easier from features based on momentum and volatility. On the other hand, the ATR demonstrates that the highest volatility stock (NVDA) is a product of its very high average stock price, which means that the volatility is expected at that high of a price. Autocorrelation determines which stocks will revert back towards the mean after any changes, and may present will more predictable trends through their history. This means that the lower the autocorrelation, the higher the mean-reverting properties, making the stock ideal for momentum based features. The larger amount of volume through the Average Daily Range and both Momentum statistics mean that any indicators involving volume may be useful to test.




Figure 1: RSI, ATR, 5-Day Momentum, and 20-Day Momentum for all seven stock tickers used in project




Figure 2: Volatility, ADR, Autocorrelation, and VP-Correlation metrics for additional feature selection support

Model Training and Prediction

The model is trained using Qiskit's Quantum Support Vector Classifier (QSVC) function which is the updated version of the QSVM function that has since been deprecated. The QSVC function requires many variables to be set up prior to running. First, a feature map must be created in order to represent the data in a quantum fashion. We must encode the classical bits that represent each feature into qubits, to do this we use qiskit's ZZFeatureMap function. This function creates a second-order Pauli-Z evolution circuit (See Fig.3) which encodes each x parameter (representing the features selected) into a qubit. This means that we can encode any amount of features selected into the same amount of qubits. The amount of qubits is directly correlated with the increase in processing time and decrease of accuracy, therefore we must be careful when navigating this issue. The solutions we determined is later discussed in the results section.




Figure 3: ZZFeatureMap Encoding Three Qubits

3 Results

After the feature map is created a qiskit StateVectorSampler is initialized that simulates the training of the QSVC classifier, this is done instead of using a real quantum computer as the project does not incorporate a high order of qubits which would require more processing power. Using the qiskit-machine-learning library we generate a FidelityQuantumKernel which is used to map the data that has now passed through the ZZFeatureMap into a higher dimensional space. This allows for the non-linearly separable data to instead be linearly separable. When some data is unable to be separated linearly, we can add in a z-axis, and then have the potential to "slice" this data through a third dimension which then allows it to be linearly separable where previously was not possible (See Fig.4).




Figure 4: Quantum Kernel classifies blue data and black data from non-linearly separable two-dimensional space (left) into linearly separable three-dimensional space (right)

Once the feature map, sampler, and quantum kernel are prepared, then the QSVC can be initialized. It uses the quantum kernel and a variable C as the

parameters. This C value represents the red line in Figure 3. A larger C value will create a larger minimum margin for the line, which means it will be closer towards the blue data, a lower c will make the red line farther towards the black data. The value of C is purely dependent on how the data looks, if there are a presence of outliers on one or the other side, then the value of C can be changed to get a better representation of all the data.

4 Evaluation

The study in which we are comparing our model tested many different models and dimensionality reduction algorithms to attempt to find the optimal combination. The highest average accuracy and f-score value they were able to achieve was 60.26% and 62.24% respectively. This was done with a classical machine learning model using a Quantum Annealing algorithm with 5 qubits. The highest QSVM metrics were achieved with the Principal Component Analysis algorithm and 5 qubits. This resulted in 56.58% accuracy and 59.79% f-score. We will be using these values as our goals to achieve. Testing was done on different hyperparameters in order to determine what would lead to the optimal prediction metrics. The interchangeable areas of the software during testing was the features, the number of repetitions of the ZZFeatureMap on the circuit, the C -value of the QSVM, and the entanglement scheme of the ZZFeatureMap.

Out of all previously mentioned hyperparameters, only one had little effect on the end outcome. The C -value of the QSVM was only changed once throughout all testing, from 10.0 to 5.0, and optimality was found with the 5.0 C -value. Therefore we determined this parameter to be insignificant when determining optimal accuracy and f-score. For the other parameters, we had tested each at different values until we reach outcomes better than the ones found in the comparative study. To test the features and their performance, we began with a baseline test over a single feature combination, this involved combining two raw data columns (previous close and volume) directly from the dataset. This was the simplest combination that would still allow the model to learn based on the stock metrics determined prior. Starting with this combination allowed us to view the general trends of each of the other hyperparameters, and changing them will affect either the runtime or the outcome of the model.




Figure 5: Baseline tests on feature combination (Previous Close, Volume) at different sample sizes, entanglement schemes, C values, and circuit repetitions

Based on the results, we determined that the effect on run-time was far more significant than the changes to the accuracy. Changes in sample size resulted in the program running anywhere from 2 minutes to 15 minutes, the effect of sample size on the runtime seemed to be exponential. We set the general testing sample size to 200 to ensure that the tests would not take too long, and so that we would get a good representation of training and testing data (160 train, 40 test). Repetitions and entanglement schemes also had similar effects on the runtime but not to the degree that the sample size had. In terms of the effect on the accuracy, none had significant effects other than C-value, which was significantly lower at 0.1 and at 100.0. This is due to the effect that the C-value has on the margin for the hyperplane, which most likely caused a significant amount of errors to occur in training. The optimal area for the C value was determined to be somewhere in between 1.0 and 10.0, we decided it to be 5.0. For repetitions and entanglement scheme, we began with using 1 and full, however this was changed later due to the circuit complexity becoming too great. Since the change in accuracy values were insignificant on repetitions and entanglement, this meant we could change them freely without much consequence.

Now that the optimal hyperparameters were determined, the most important task of feature selection was needed to be done. This required testing multiple different possible feature combinations until we found a combination that outperformed the highest average accuracy and f-score of the study we are

comparing with. There are two measures that are important for feature selection; number of features (qubits) and the type of combination. In order to achieve an outperforming prediction model we needed to find the proper amount of qubits so that the circuit was not too complex and thus result in excessive noise, as well as find the right combination of features to ensure we are covering every type of stock trend. From the metrics in figure 1 and 2 we determined that volume, momentum, volatility are all important markers for these stocks, these will be the key factors to test in order to help predict the trends optimally. The first set of tests were done in order to correlate the highest outcomes with specific indicators, if one particular indicator was linked to multiple instances of high performance, then that indicator was determined to be ideal.

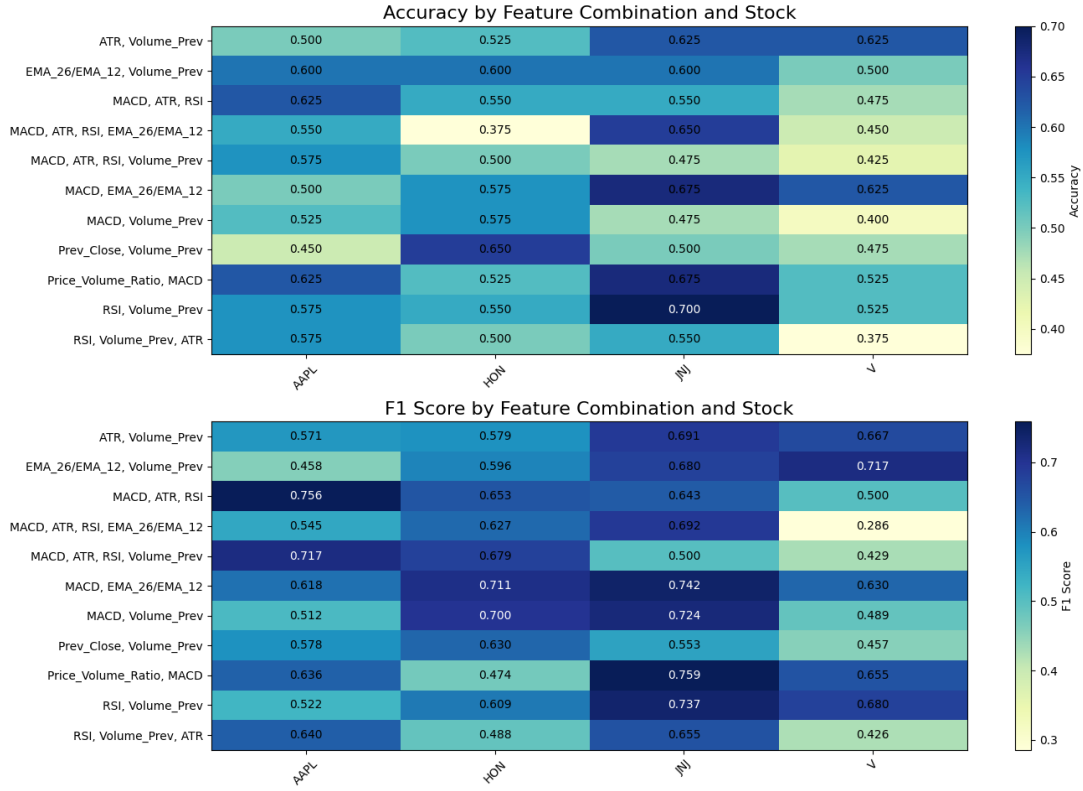


Figure 6: Initial tests on correlation between featrure combinations and (f1 score, accuracy)

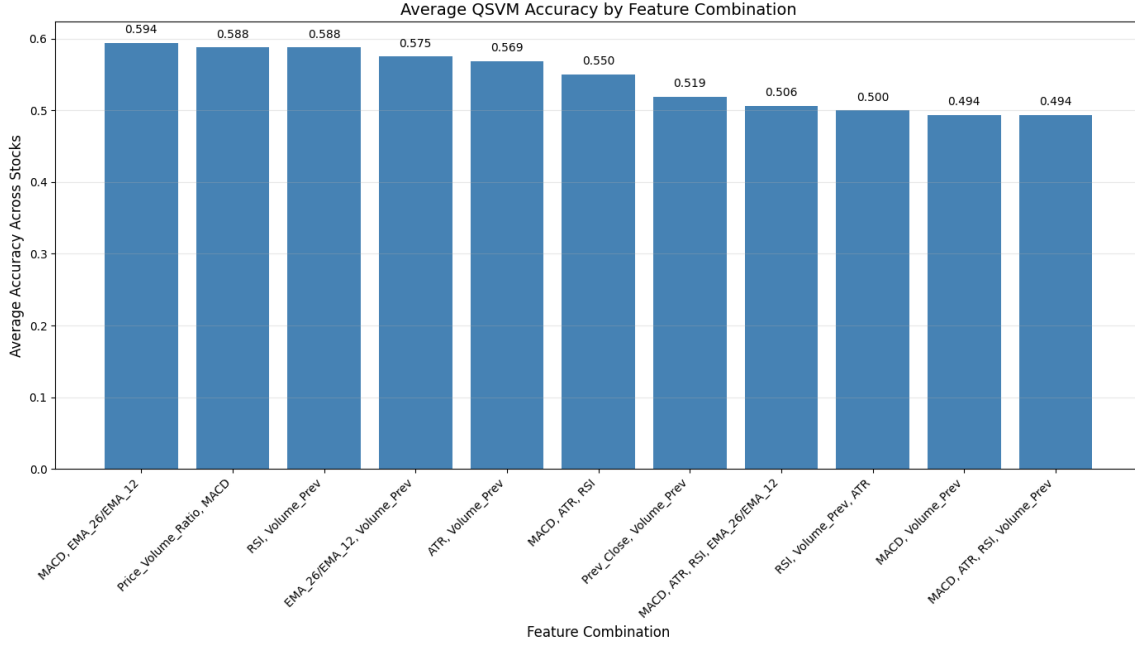


Figure 7: Average accuracy of different feature combinations throughout all stock predictions

The first tests helped determine the indicators that seemed to perform better than others. We found that feature combinations of more than 2 resulted in lower average accuracy (See figure 7), therefore it was optimal to pick just two features to ensure accuracy remained highest. The choices of features were based on utilizing more stock indicators instead of raw data, in this way we could determine which of the specific indicators has more of a presence in high performing QSVM predictions. The first tests also proved to be promising as we noticed that most feature combination choices at minimum replicated withing 3-5% of the average scores found in the paper. The highest performing feature combination of (MACD) and the EMA Ratio (EMA26/EMA12), was able to outperform the average average accuracy and average f1-score of the highest QSVM implemented in the paper. The QSVM with Principal Component Analysis and 5 features achieved an accuracy of 56.58% and f-score of 59.79% while our implementation of the QSVM achieved an accuracy of 59.4% and f-score of 67.5%.

The difference in sectors among the stocks chosen provided the QSVM with learning data that has wildly different patterns. This was demonstrated in the testing when looking at the performance of feature combinations for specific stocks. With the chosen feature combinations, a C-value of 5.0, 2 repetitions of the ZZFeatureMap, and linear entanglement, we were able to produce a

QSVM that outperformed the paper on every metric (except for HON f-score) for each different stock (See figure 8). These difference was much greater than the difference present in the average accuracy and f1-score value. They do not hold as much weight as the total average accuracy and f1-score values do, nonetheless it is still very notable for this project.

AAPL							
Best Features	Project Accuracy	Project F1 Score	Paper Method	Paper Accuracy	Paper F1 Score	Accuracy Difference	F1 Score Difference
			Quantum SVM Circular Quantum				
MACD,ATR,RSI	63%	76%	Annealing-5	59%	58%	4%	18%
HON							
Best Features	Project Accuracy	Project F1 Score	Paper Method	Paper Accuracy	Paper F1 Score	Accuracy Difference	F1 Score Difference
Previous Close, Previous Volume	65%	63%	Quantum SVM Full/Linear PCA-5	58%	66%	7%	-3%
JNJ							
Best Features	Project Accuracy	Project F1 Score	Paper Method	Paper Accuracy	Paper F1 Score	Accuracy Difference	F1 Score Difference
			Quantum SVM Circular Quantum				
RSI, Previous Volume	70%	74%	Annealing-5	60%	57%	10%	17%
V							
Best Features	Project Accuracy	Project F1 Score	Paper Method	Paper Accuracy	Paper F1 Score	Accuracy Difference	F1 Score Difference
ATR, Previous Volume	63%	67%	Quantum SVM Full/Circular PCA-3	57%	62%	6%	5%

Figure 8: Table demonstrating comparison of our QSVM with specific feature combinations vs. paper's QSVM model on predicting specific stocks

5 Conclusion

5.1 Summary

The complexity of predicting stock prices using quantum algorithms and methods remains a challenge. The overwhelming unpredictability that the stock market possesses makes such a task very difficult. The results from this project aim to give details on methods to further optimize the usage of QSVM's in such a binary classification problem. This project revealed the grave importance of testing different parameters in order to develop further optimizations, especially related to the usage of features in the QSVM.

5.2 Relevance

The content of this project is particularly relevant to the course as it involves heavily with developed further optimized structure in a quantum fashion to achieve a goal. The ability to transform a difficult classical issue into a easier quantum one has been the forefront of the course since it began. The optimization towards stock price prediction sheds light on the both the ability that quantum computers have in binary classification, and the potential for greater limits in the field.

5.3 Future Work

For this project in particular, many improvements can be made in order to develop this into higher optimality. Development of a feature selection algorithm that allows for dimensionality reduction while maintaining the mathematically most optimal features is something that could be incorporated in the future. The inclusion of real quantum computers is another topic that could be developed in the future. Utilizing IBM's real quantum computer could have the processing capabilities to predict stock prices with many more qubits that we were able to simulate, leading to further developments towards a better model.

Contribution of Team Members

Code Implementation

- A. Benjamin Muir implemented the Feature Map, main QSVM functionality, and data testing variant of Code
- B. Hasan Fakih implemented the stock indicator helper functions, model training and testing
- C. Both members worked equally on the data pre-processing/preparation of data

Report

- A. Benjamin Muir worked on Results, Evaluation and Conclusion
- B. Hasan Fakih worked on Introduction and Background Information
- C. Benjamin Muir and Hasan Fakih both jointly worked on the research and finding papers relevant to the project

References

- [1] Xia, S., Wei, M., Zhu, Y., & Pu, Y. (2024). AI-Driven Intelligent Financial Analysis: Enhancing Accuracy and Efficiency in Financial Decision-Making. *Journal of Economic Theory and Business Management*, 1(5), 1–11. <https://doi.org/10.5281/zenodo.13926298>
- [2] Parshv Chhajer, Manan Shah, Ameya Kshirsagar, The applications of artificial neural networks, support vector machines, and long–short term memory for stock market prediction, *Decision Analytics Journal*, Volume 2, 2022, 100015, ISSN 2772-6622, <https://doi.org/10.1016/j.dajour.2021.100015>
- [3] Alexander, A., & Widdows, D. (2022). Quantum text encoding for classification tasks. 2022 IEEE/ACM 7th Symposium on Edge Computing (SEC), 355 – 361. <https://doi.org/10.1109/sec54971.2022.00052>
- [4] Bhasin, N. K., Kadyan, S., Santosh, K., HP, R., Changala, R., & Bala, B. K. (2024). Enhancing quantum machine learning algorithms for optimized financial portfolio management. 2024 Third International Conference on Intelligent Techniques in Control, Optimization and Signal Processing (INCOS), 1 – 7. <https://doi.org/10.1109/incos59338.2024.10527612>
- [5] Onyshchak, O. (2020, April 2). Stock market dataset. Kaggle. <https://www.kaggle.com/datasets/jacksoncrow/stock-market-dataset/data>
- [6] Srivastava, N., Belekhar, G., Shahakar, N., & Babu H., A. (2023). The potential of quantum techniques for stock price prediction. 2023 IEEE International Conference on Recent Advances in Systems Science and Engineering (RASSE), 1 – 7. <https://doi.org/10.1109/rasse60029.2023.10363533>

Acronyms

AI Artificial Intelligence. 1, 15

ANN Artificial Neural Network. 2, 15

ATR Average True Range. 3, 4, 15

EMA Exponential Moving Average. 2, 3, 11, 15

LSTM Long Short-Term Memory. 2, 15

MACD Moving Average Convergence Divergence. 2, 3, 11, 15

ML Machine Learning. 1, 2, 15

PCA Principal Component Analysis. 2, 3, 15

QML Quantum Machine Learning. 2, 15

QSVC Quantum Support Vector Classifier. 6, 7, 15

QSVM Quantum Support Vector Machine. 1–4, 6, 8, 11–13, 15

QUBO Quadratic Unconstrained Binary Optimization. 4, 15

RSI Relative Strength Index. 2, 3, 15

SVM Support Vector Machine. 1, 2, 15